

Signals (POSIX systems)



Patrício R. Domingues

Departamento de Eng Informática

ESTG/IPLeiria

POINTER TO FUNCTIONS

Pointer to functions (1)

- The C language has the concept of *pointer to a function*
 - The pointer holds the address of a function
 - The function can be called through the pointer
 - Call syntax is similar to a regular function call
- What's the datatype of a pointer to a function?
 - Datatype is function
 - A function has a signature defined by the following elements
 - Datatype of the return value
 - Datatype of the parameters (if any)

Pointer to functions (2)

- Examples of signatures of functions
 - `int F1(int a, int b);`
 - `int F2(int y, int z);`
 - F1 and F2 have the same signature (the name of the parameters is irrelevant)
 - `double F3 (int a, int b);`
 - F3 has a signature different from F1 (and obviously from F2)
 - `int *F4(int a, int b);`
 - F4 has a signature different from F1, F2 and F3

Pointer to functions (3)

- Declaration of a pointer to a function
 - Pointer points to the signature
- Example
 - `int (*PtrF1)(int , int);`
 - PtrF1 can point to functions with signature “**int ... (int, int);**”
 - `double (*PtrF2)(double ,char *);`
 - PtrF2 can point to functions with signature “**double ... (double,char*);**”

Pointer to functions (4)

- How to get the address of a function?
 - simple
 - The name of the function holds the address of the function
 - Thus, the pointer only needs to be assigned the name of the function
 - This can also be done through “**&FunctionName**”

- Example

```
int F1(int a, int b){  
    return a + b;  
}
```

```
int (*PtrF1)(int , int) = NULL; /* PtrF1 declaration */
```

```
int Result;
```

```
PtrF1 = F1; /* PtrF1 points to the F1 function */
```


```
Result = PtrF1(10,30); /* Call of F1 through PtrF1 */
```

```
Result = (*PtrF1)(10,30); /* Same as previous line of code */
```

SIGNALS

What's a *signal*? (1)

- Signal is an asynchronous notification delivered to a process
 - Notification
 - Related to an event
 - Example:
 - » control + C within the shell has the effect of delivering a signal (SIGINT) to the running process
 - » Process tries to access an invalid memory address. SO delivers the SIGSEGV (segmentation violation) to the process.
 - Asynchronous
 - The signal can happen at any time
 - Therefore, the notification can be delivered at any time...



<https://bit.ly/2wLTTca>

What's a *signal*? (2)

- In UNIX, a signal is represented by an integer number
- Each signal has also a symbolic name
 - Easier for us humans to remember
 - Examples: SIGINT, SIGHUP, SIGSEGV, SIGPIPE, etc.
- There are around 40 different signals
 - man 7 signal
 - kill -l signal
 - List the available signals

signals

drawings.juns.ca

If you've ever used
kill
you've used signals



the Linux kernel sends your process signals in lots of situations



Every signal has a default action, one of:

ignore


kill process

kill process AND make core dump file

stop process

resume process

Your program can set custom handlers for almost any signal



exceptions:

SIGSTOP & SIGKILL
can't be ignored
got →
SIGKILLED

you can send signals yourself with the **kill** system call or command

SIGINT ctrl-C } various levels of
SIGTERM kill } "die"
SIGKILL kill -9 }

SIGHUP kill -HUP

↑
often interpreted as
"reload config", eg by nginx

Signals can be hard to handle correctly since they can happen at ANY time

handling a signal

SURPRISE!
another signal!



Sources of signals (#1)

- Signals come from many sources. Some possible sources:
 - Something executing the system call (“`kill(...)`”) that transmit a signal to a process
 - Sending a signal from the process unto itself
 - When a child process exits
 - When the parent process dies or *hangup*
 - When the program behaves incorrectly
 - Hardware failure

Source of signals

- Which elements can send a signal to a process?




Figure: Dr Eileen T. Kramer
(University of Georgia, USA)

The *kill* command (1)

- kill command (tries to) deliver a signal to a destination process
 - Syntax
 - kill [options] PID1 PID2 PID3...
 - Send *signal* to processes with PID1, PID2, PID3, ...
 - Why the name *kill*?
 - By default, when a process receives a signal, it terminates
 - A process can be configured to have a different behavior
 - *Handling* the signal
 - More on the command kill
 - man kill

The *kill* command (2)

- Examples

- `kill -SIGINT 1234`

- Send the signal SIGINT to process with PID 1234

- `kill -9 1234`

- Send the signal 9 (**SIGKILL**) to processes 1234

- `kill -SIGKILL 1234`

- Same as above

- `kill -kill $$`

- Kills the shell
 - \$\$ is the PID of the shell

SIGKILL and SIGSTOP are the only signal that cannot be captured (more on this later)

Sending a signal

- A regular user can only send a signal to processes that she/he owns
 - kill -KILL 1

```
bash: kill: (1) - Operation not permitted
```
 - Process with PID 1 is a system process (process *init*)
- A root account can send a signal to any process

The *killall* command

- Killall is a variant of the kill command
- `killall [options] name`
 - It sends the signal to all processes whose name is *name*
- `killall -INT bash`
 - (*tries*) to send the *SIGINT* signal to all processes whose name is *bash*
 - For a regular user, only the bash processes that belong to the user receive the signal



IPL

escola superior de tecnologia e gestão
instituto politécnico de leiria

kill doesn't just kill programs



you can send ANY signal to a program with kill!

kill - SIGNAL PID

name or number

killall - SIGNAL NAME

signals all processes called NAME for example

\$ killall firefox

useful flags:

-w wait for all signaled processes to die

-i ask before signalling

kill

which signal kill sends

| | <u>name</u> | <u>num</u> |
|------------|-------------|------------|
| kill | => SIGTERM | 15 |
| kill -9 | => SIGKILL | 9 |
| kill -KILL | | |
| kill -HUP | => SIGHUP | 1 |
| kill -STOP | => SIGSTOP | 19 |

JULIA EVANS
@b0rk

kill -l

lists all signals

| | | | |
|---------|-----------|---------|-----------|
| 1 HUP | 2 INT | 3 QUIT | 4 ILL |
| 5 TRAP | 6 ABRT | 7 BUS | 8 FPE |
| 9 KILL | 10 USR1 | 11 SEGV | 12 USR2 |
| 13 PIPE | 14 ALRM | 15 TERM | 16 STKFLT |
| 17 CHLD | 18 CONT | 19 STOP | 20 TSTP |
| 21 TTIN | 22 TTOU | 23 URG | 24 XCPU |
| 25 XFSZ | 26 VTALRM | 27 PROF | 28 WINCH |
| 29 POLL | 30 PWR | 31 SYS | |

pgrep

prints PIDs of matching running programs

pgrep fire matches firefox, firebird
NOT bash firefox.sh

To search the whole command line (eg bash firefox.sh)
use pgrep -f

pkill

same as pgrep , but signals PIDs found. ex:

pkill -9 -f firefox



I use pkill more than killall these days

The kill() function (#1)

- How do we send a signal programmatically?
 - kill function
 - `int kill(pid_t pid, int sig);`
 - return -1 on error (setting *errno*), 0 on success
 - Interpretation of pid
 - $pid > 0$: send signal to process with PID *pid*
 - $pid == 0$: send signal to all processes whose group ID equals the sender
 - » Parent process sending a signal to all children
 - Interpretation of sig
 - $Sig > 0$: send signal *sig*
 - $Sig == 0$: do not actually send a signal, but acts if has done so
 - » Use-case: checking if a given process exist

The kill() function (#2)

- man 2 kill
 - Need to specify the section 2 of the manual
 - Section 2 of man describes system calls
 - Kill, read, write, open...
 - man 2 intro

INTRO(2)

Linux Programmer's Manual

INTRO(2)

NAME

intro - introduction to system calls

DESCRIPTION

Section 2 of the manual describes the Linux system calls. A system call is an entry point into the Linux kernel. Usually, system calls are not invoked directly: instead, most system calls have corresponding C library wrapper functions which perform the steps required (e.g., trapping to kernel mode) in order to invoke the system call. Thus, making a system call looks the same as invoking a normal library function.

For a list of the Linux system calls, see **syscalls(2)**.

Signals sent by functions

- Some system functions send a signal
 - **abort()** sends the SIGABRT signal to the calling process
 - `void abort(void);`
 - **alarm()** schedule the delivery of the SIGALRM signal to the calling process
 - Appropriate for setting timeouts
 - `unsigned int alarm(unsigned int seconds);`

Responding to a signal

- For given signal, a process can be configured to...
 - Execute the default action
 - The default action for many (not all!) signals is to terminate the process
 - SIGKILL and SIGSTOP always results in their default actions
 - SIGKILL: terminates the process
 - SIGSTOP: stops the process (same as ctrl+Z on the shell)
 - Ignore the signal
 - Not available for SIGKILL and SIGSTOP
 - Launch a signal *handler*
 - *Signal handler*: function called whenever a signal is received

Setting a *signal handler*

- The **sigaction** library call
 - Installs a signal handler for a given signal

```
#include <signal.h>

int sigaction(int signum, const struct
              sigaction *act, struct sigaction *oldact);
```

- **signum**: signal number to be configured
- **act**: pointer to **sigaction** struct that contains the configuration to use
- **oldact**: **sigaction** struct filled with the previous configuration
- Question
 - Why const **struct sigaction *act** **vs.** **struct sigaction *oldact**?

struct sigaction (#1)

- The *struct sigaction*

```
struct sigaction {  
    void (*sa_handler) (int);  
    void (*sa_sigaction) (int, siginfo_t *, void *);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_restorer) (void); /* << obsolete */  
};
```

sa_handler:

- pointer to the function that will handle the signal. The function should have one integer parameter and does not return anything
 - The integer parameter corresponds to the signal number that triggered the handler
- SIG_DFL to restore the default behavior
- SIG_IGN to set the process to ignore the signal

struct sigaction (#2)

- The *struct sigaction*

```
struct sigaction {  
    void (*sa_handler)(int);  
    void (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_restorer)(void); /* << obsolete */  
};
```

sa_sigaction:

- If `sa_flags` is `SA_SIGINFO`, `sa_sigaction` specifies the signal handling function
- This has improved functionalities for the signal handler function, namely a `siginfo_t` structure (see next slide)
- `sa_sigaction` is not compatible with `sa_handler`. Use one or the other, but not both!

struct sigaction (#3)

- The *struct siginfo_t*

```
struct siginfo_t {  
    int si_signo;      /* Signal number */  
    int si_errno;      /* An errno value */  
    int si_code;       /* Signal code */  
    pid_t si_pid;      /* Sending process ID */  
    uid_t si_uid;      /* Real user ID of sending process */  
    int   si_status;    /* Exit value or signal */  
    clock_t si_utime;  /* User time consumed */  
    clock_t si_stime;  /* System time consumed */  
(...)  
}
```

- The `struct siginfo_t` returns a large number of data regarding the signal and the context of the call

struct sigaction (#4)

- The *struct sigaction*

```
struct sigaction {  
    void (*sa_handler)(int);  
    void (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_restorer)(void); /* << obsolete */  
};
```

sa_mask:

- Set of signal to be blocked during the call of a signal handler
 - Avoid race conditions
 - During the handler execution, the signal being processed by the handler is blocked by default
 - Regarding `sigset_t`, see `man 3 sigsetops`

struct sigaction (#5)

- The *struct sigaction*

```
struct sigaction {  
    void (*sa_handler) (int);  
    void (*sa_sigaction) (int, siginfo_t *, void *);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_restorer) (void); /* << obsolete */  
};
```

sa_flags:

- specifies a set of flags which modify the behavior of the signal.
- It can comprise several flags (through bitwise OR)
- Example
 - SA_NOCLDSTOP | SA_ONSTACK | ...
 - SA_SIGINFO to use the field `sa_sigaction` as handler
 - | is the bitwise OR operator
 - See **man 3 sigaction**

Example with *sigaction*

```
int main(void)  {
    struct sigaction act;
    act.sa_handler = process_signal;
    sigemptyset( &act.sa_mask );
    act.sa_flags = 0;
    sigaction( SIGINT, &act, 0 );
    while(1)
    {
        printf("waiting one second - press CTRL+C :)\n");
        sleep(1);
    }
    return 0;
}

void process_signal(int signum){
    printf("Capturing %d (SIGINT=%d)\n", signum, SIGINT);
}
```

The program will not terminate with
CTRL+C (SIGINT)
However, it will terminate with CTRL+\
(SIGQUIT)

The printf function is not an async-signal-safe (it can lead to race conditions)
The “write” function is async-signal-safe (see man 7 signal-safety)

- SIGINT can be sent to a foreground process with CTRL+C



Example with siginfo_t (1)

```
/* Setting sigaction() to use siginfo_t. */
static void hdl(int sig, siginfo_t *siginfo, void *context) {
    printf ("Sending PID: %ld, UID: %ld\n",
    (long)siginfo->si_pid, (long)siginfo->si_uid);
}

int main (int argc, char *argv[]) {
    struct sigaction act;
    memset (&act, '\0', sizeof(act));
    /* Use sa_sigaction field: handle has
two additional parameters */
    act.sa_sigaction = &hdl;
```

Example with siginfo_t (2)

```
/* SA_SIGINFO flag tells sigaction() to use the sa_sigaction
field, not sa_handler. */

act.sa_flags = SA_SIGINFO;

if (sigaction(SIGTERM, &act, NULL) < 0) {
    perror ("sigaction");
    return 1;
}

while (1) {
    sleep(10);
}

return 0;
}
```

System calls and signals

- When a system call (e.g. `read()`) is interrupted by a signal...
 - 1) The signal handler is called
 - 2) The signal handler terminates and thus returns the control back to the system call
 - On UNIX, slow system calls do not resume.
 - Whenever a signal is received within a system call, the system call returns an error and sets `errno` to `EINTR`
- What is a *slow system call*?
 - System calls that perform I/O operations on devices that can block the caller *forever*
 - sockets (networks), pipes
 - The `pause()` system call
 - Blocks the process until it receives a signal
- Use `SA_RESTART` to recover slow system calls automatically

pause function

- The pause(2) function
 - Suspends execution until any signal is caught.

PAUSE(2) Linux Programmer's Manual PAUSE(2)

NAME
pause - wait for signal

SYNOPSIS

```
#include <unistd.h>

int pause(void);
```

DESCRIPTION

pause() causes the calling process (or thread) to sleep until a signal is delivered that either terminates the process or causes the invocation of a signal-catching function.

RETURN VALUE

pause() returns only when a signal was caught and the signal-catching function returned. In this case, pause() returns -1, and errno is set to EINTR.

ERRORS

EINTR a signal was caught and the signal-catching function returned .

CONFORMING TO


POSIX.1-2001, POSIX.1-2008, SVr4, 4.3BSD.

Manual page pause(2) line 1 (press h for help or q to quit)

PROCESS GROUPS, SESSIONS & SIGNALS

Process group

- In addition to the parent/child process association, there are two more associations:
 - Process group: PGID
 - Sessions: SID
- bash creates a process group for each command line that has pipes
- Each process group has a unique PGID
- Command "ps j" lists the processes with the indication of the PGID and SID
- Example: **tail -f /var/log/syslog | grep "CRON" &**




The screenshot shows a terminal window with the following content:

```
osboxes@osboxes:~$ tail -f /var/log/syslog | grep "CRON" &
[1] 2420
osboxes@osboxes:~$ ps j
PPID  PID  PGID   SID TTY      TPGID STAT  UID   TIME COMMAND
 1869 1872 1872 1872 pts/0    1872 Ss+  1000  0:00 /bin/bash
 1869 2057 2057 2057 pts/1    2057 Ss+  1000  0:00 /bin/bash
 2057 2395 2395 2057 pts/1    2057 S1   1000  0:00 /usr/bin/python3
 2395 2402 2402 2402 pts/2    2421 Ss   1000  0:00 /bin/bash
 2402 2419 2419 2402 pts/2    2421 S    1000  0:00 tail -f /var/log
 2402 2420 2419 2402 pts/2    2421 S    1000  0:00 grep --color=aut
 2402 2421 2421 2402 pts/2    2421 R+   1000  0:00 ps j
```

A red box highlights the "PGID" column header in the ps j command output. Two red arrows point from the "PGID" header to the "PGID" value of 2402, which appears in the first three rows of the output.

Sessions (SID)

- Session is a Process group set
 - Usually associated with:
 - A control terminal
 - A process designated as a session leader
 - If a session has a control terminal, it will have a single foreground process group, and all other process groups in the session will be background groups
 - Example: BASH with PID 2402 is session leader SID=2402 which also contains tail (PID=2419) and grep (PID=2420)



The screenshot shows a terminal window titled "osboxes@osboxes: ~". The title bar also indicates the window ID is 83x24. The terminal displays two command lines:

```
osboxes@osboxes:~$ tail -f /var/log/syslog | grep "CRON" &
[1] 2420
osboxes@osboxes:~$ ps j
```

Following these commands, the terminal lists process information in a table format:

| PPID | PID | PGID | SID | TTY | TPGID | STAT | UID | TIME | COMMAND |
|------|------|------|------|-------|-------|------|------|------|------------------|
| 1869 | 1872 | 1872 | 1872 | pts/0 | 1872 | Ss+ | 1000 | 0:00 | /bin/bash |
| 1869 | 2057 | 2057 | 2057 | pts/1 | 2057 | Ss+ | 1000 | 0:00 | /bin/bash |
| 2057 | 2395 | 2395 | 2057 | pts/1 | 2057 | S1 | 1000 | 0:00 | /usr/bin/python3 |
| 2395 | 2402 | 2402 | 2402 | pts/2 | 2421 | Ss | 1000 | 0:00 | /bin/bash |
| 2402 | 2419 | 2419 | 2402 | pts/2 | 2421 | S | 1000 | 0:00 | tail -f /var/log |
| 2402 | 2420 | 2419 | 2402 | pts/2 | 2421 | S | 1000 | 0:00 | grep --color=aut |
| 2402 | 2421 | 2421 | 2402 | pts/2 | 2421 | R+ | 1000 | 0:00 | ps j |

The columns are labeled: PPID, PID, PGID, SID, TTY, TPGID, STAT, UID, TIME, and COMMAND. The "SID" column is highlighted with a red box.

Signals/ GPID / SID


- Sending a signal to a process group

- Kill command
 - Syntax: **kill -SIGTERM -PGID**
 - Example: **kill -SIGTERM -2402**

Indicate minus sign

- Sending a signal to a SID session

- **pkill** Command
 - Syntax: **pkill -s SID**
 - Example: **pkill -s 2402**



The screenshot shows a terminal window titled "osboxes@osboxes: ~". The user has run the command "tail -f /var/log/syslog | grep \"CRON\" &". Then, they ran "ps j" to list processes. A red box highlights the "SID" column header. The output shows several processes, with the last three rows being the grep command itself:

| PPID | PID | PGID | SID | TTY | TPGID | STAT | UID | TIME | COMMAND |
|------|------|------|------|-------|-------|------|------|------|-------------------|
| 1869 | 1872 | 1872 | 1872 | pts/0 | 1872 | Ss+ | 1000 | 0:00 | /bin/bash |
| 1869 | 2057 | 2057 | 2057 | pts/1 | 2057 | Ss+ | 1000 | 0:00 | /bin/bash |
| 2057 | 2395 | 2395 | 2057 | pts/1 | 2057 | S1 | 1000 | 0:00 | /usr/bin/python3 |
| 2395 | 2402 | 2402 | 2402 | pts/2 | 2421 | Ss | 1000 | 0:00 | /bin/bash |
| 2402 | 2419 | 2419 | 2402 | pts/2 | 2421 | S | 1000 | 0:00 | tail -f /var/log |
| 2402 | 2420 | 2419 | 2402 | pts/2 | 2421 | S | 1000 | 0:00 | grep --color=auto |
| 2402 | 2421 | 2421 | 2402 | pts/2 | 2421 | R+ | 1000 | 0:00 | ps j |

SIGNALS IN LINUX

Signals in Linux (#1)

- Source: “*Signals*”, *Chapter 10 - Linux System Programming*, Robert Love, 2nd Edition, O'Reilly, 2013

| Signal | Description | Default action |
|---------|--|--------------------------|
| SIGABRT | Sent by <code>abort()</code> | Terminate with core dump |
| SIGALRM | Sent by <code>alarm()</code> | Terminate |
| SIGBUS | Hardware or alignment error | Terminate with core dump |
| SIGCHLD | Child has terminated | Ignored |
| SIGCONT | Process has continued after being stopped | Ignored |
| SIGFPE | Arithmetic exception | Terminate with core dump |
| SIGHUP | Process's controlling terminal was closed (most frequently, the user logged out) | Terminate |
| SIGILL | Process tried to execute an illegal instruction | Terminate with core dump |
| SIGINT | User generated the interrupt character (Ctrl-C) | Terminate |
| SIGIO | Asynchronous I/O event | Terminate ^a |
| SIGKILL | Uncatchable process termination | Terminate |
| SIGPIPE | Process wrote to a pipe but there are no readers | Terminate |
| SIGPROF | Profiling timer expired | Terminate |

Signals in Linux (#2)

- Source: “*Signals*”, *Chapter 10 - Linux System Programming*, Robert Love, 2nd Edition, O'Reilly, 2013

| Signal | Description | Default action |
|-----------|---|--------------------------|
| SIGSEGV | Memory access violation | Terminate with core dump |
| SIGSTKFLT | Coprocessor stack fault | Terminate ^b |
| SIGSTOP | Suspends execution of the process | Stop |
| SIGSYS | Process tried to execute an invalid system call | Terminate with core dump |
| SIGTERM | Catchable process termination | Terminate |
| SIGTRAP | Break point encountered | Terminate with core dump |
| SIGTSTP | User generated the suspend character (Ctrl-Z) | Stop |
| SIGTTIN | Background process read from controlling terminal | Stop |
| SIGTTOU | Background process wrote to controlling terminal | Stop |
| SIGURG | Urgent I/O pending | Ignored |
| SIGUSR1 | Process-defined signal | Terminate |
| SIGUSR2 | Process-defined signal | Terminate |
| SIGVTALRM | Generated by <code>setitimer()</code> when called with the <code>ITIMER_VIRTUAL</code> flag | Terminate |
| SIGWINCH | Size of controlling terminal window changed | Ignored |
| SIGXCPU | Processor resource limits were exceeded | Terminate with core dump |
| SIGXFSZ | File resource limits were exceeded | Terminate with core dump |

^aThe behavior on other Unix systems, such as BSD, is to ignore this signal.

^bThe Linux kernel no longer generates this signal; it remains only for backward compatibility.

GETTING A SIGNAL'S NAME

Signal name in Linux

- The system-defined `sys_siglist[...]` vector of strings holds names of each signal
 - `extern const char *const sys_siglist[];`
- The name of signal is also available with:
 - `char *strsignal(int sig);`
- Example:

```
#include <signal.h>
for(i=0;i<20;i++) {
    printf("signal %d => '%s'\n", i, sys_siglist[i]);
}
```

References

- “*Signals*”, *Chapter 10 - Linux System Programming*, Robert Love, 2nd Edition, O'Reilly, 2013
- “UNIX SIGNALS”, Overview of signals
 - <https://venam.nixers.net/blog/unix/2016/10/21/unix-signals.html>
 - Audio: <https://nixers.net/showthread.php?tid=2003>
- Sigaction
 - <http://www.linuxprogrammingblog.com/code-examples/sigaction>

